



# Neural Architecture Search

Sara Boyd, Lauren Gillespie, Elyssa Sliheet



# Motivation

- DNN meta-architecture drives network success
- Hand-designed architectures are successful in certain domains [5, 6]
- How to find good architecture for arbitrary domain?
- Enter **neural architecture search (NAS)**
- NAS: finding a good network architecture automatically [7]
  - Semantics, sometimes includes hyperparameter search, sometimes separate
- Our goal: NAS using a simple EA



# DNN- Design Choices

1. Number of layers
2. Type - convolutional, pooling, fully connected
3. The ordering of layers
4. Hyperparameters for each type of layer
  - a. Receptive field size
  - b. Stride
  - c. Number of receptive fields for convolution layer
  - d. Activation function
  - e. Dropout rate

The number of possible choices makes the design space of CNN architectures extremely large and hence, infeasible for an exhaustive manual search.

# Design the CNN using Reinforcement Learning [1]

Uses a novel Q-learning agent


Goal: discover CNN architectures that perform well on a given ML task with no human intervention

Agent has a large **space of model architectures** to search from

Learns through random exploration

$\epsilon$ -greedy

Validation accuracy = reward



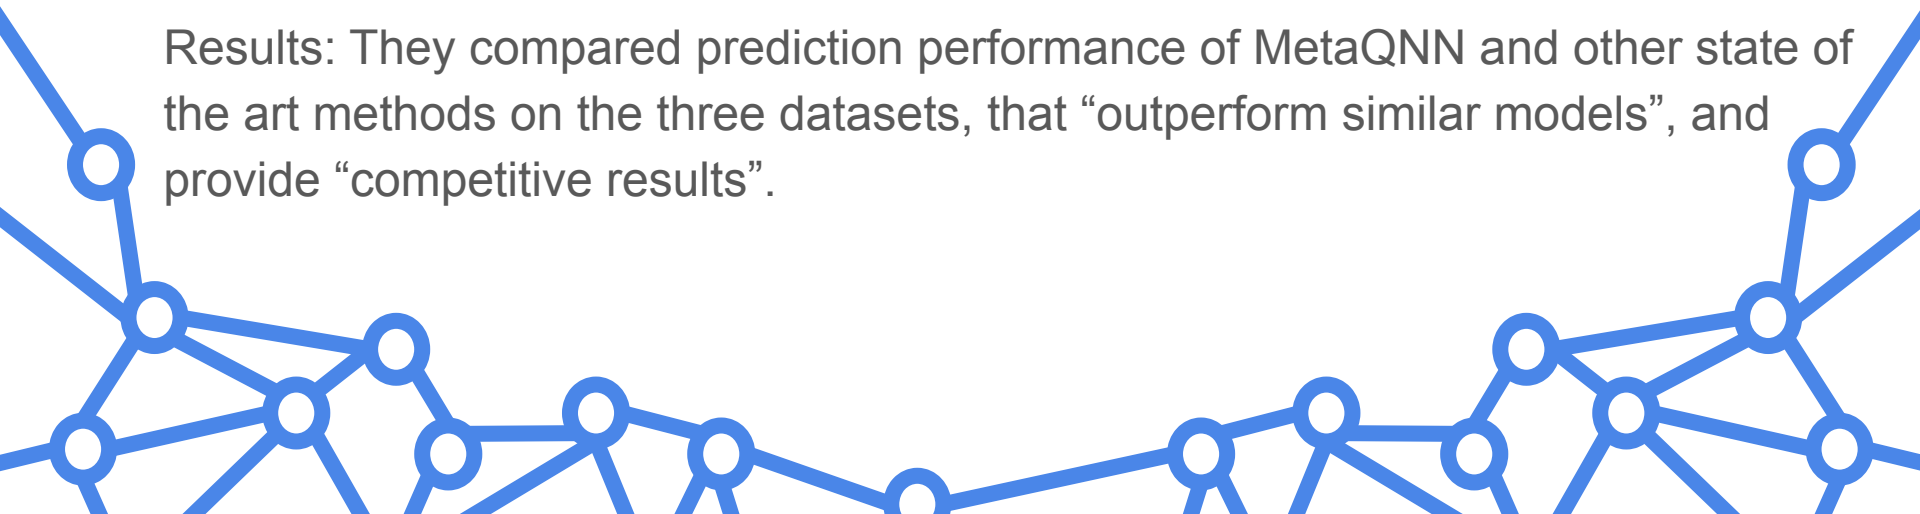
# Design the CNN using Reinforcement Learning [1]

MetaQNN - Q-learning based meta modeling

Three standard image classification datasets: CIFAR-10 [2], MNIST [6], SVHN [3]


Task: train an agent to sequentially choose neural network layers

Results: They compared prediction performance of MetaQNN and other state of the art methods on the three datasets, that “outperform similar models”, and provide “competitive results”.




# Evolutionary algorithm hyperparameter optimization

- Evolutionary algorithms can do Neural Architecture Search [7,8]
- CoDeepNEAT: NEAT for DNNs - nodes represent whole layers [7]
  - State-of-the-art on Wikipedia Comment Toxicity & Chest X-rays Multitask Image Classification
  - Architecture search better than hyperparameter search
- Simple EA on CIFAR-10 and CIFAR-100 dataset [8]
  - achieved 94.6% and 77% respectively
  - EA with tournament selection
  - Relevant parameters: learning rate, weights, convolutional layer insertion, convolutional layer deletion, stride, channel number, filter size
- Limitations of these systems: computational & time cost




# Our Initial Attempt

- Manually modified the network for training on the Keras IMDB dataset
- Number of layers
  - Activation function
    - relu, sigmoid, tanh, softmax, elu, selu, softplus, softsign, exponential, linear
- Drawbacks: manual manipulation, architecture already optimized



# Initial Attempt Results



```
model = keras.Sequential()  
model.add(keras.layers.Embedding(vocab_size, 16))  
model.add(keras.layers.GlobalAveragePooling1D())  
model.add(keras.layers.Dense(16, activation=tf.nn.relu))  
model.add(keras.layers.Dense(1, activation=tf.nn.sigmoid))
```




# Changed Direction

- Develop a program to search for an optimal solution using an  $n+1$  strategy
- Activation functions
  - tanh, softmax, elu, selu, softplus, softsign, relu, sigmoid, hard\_sigmoid, exponential, linear
- Network structure - all layers are dense
  - Input
  - Activation (modified parameter)
  - Dropout (modified parameter)
  - Dense layer with num\_classes param
  - Activation (modified parameter)
- Child generation
  - Keep the network with the highest accuracy rate (elite population = 1)
  - Fill the rest of the generation with networks that have two random activation functions and a random dropout rate (0-1 range)
- Datasets
  - Reuters Text Categorization Dataset

# Results


Epoch: 5, Population size: 5, Number of Generations: 20

REUTERS DATASET	Input Activation	Dropout Rate	Output Activation	Accuracy
Default Architecture	relu	0.5	softmax	0.79
Our Default Architecture	linear	0.0	linear	0.38
Our Architecture	sigmoid	0.08	sigmoid	0.80



# Future Work

- Testing with MNIST dataset, but not fully functional
- Include modifying the number of layers (neural architecture search)
- Increase population size and elitism parameter/tournament selection
- Parallelize to use physics department NVIDIA GPUs
- Add more parameters to modify (kernel size, padding, optimizers)
- Modify shape of layers





# References

- [1] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, November 1998.
- [2] Karen Simonyan\* & Andrew Zisserman. VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION. In *Proceedings of International Conference on Reinforcement Learning 2015*.
- [3] Jason Liang, Elliot Meyerson, Babak Hodjat, Dan Fink, Karl Mutch, Risto Miikkulainen. Evolutionary Neural AutoML for Deep Learning. In *Proceedings of Conference on Genetic and Evolutionary Computation 2019*.
- [4] Bowen Baker, Otkrist Gupta, Nikhil Naik, Ramesh Raskar. Designing Neural Network Architectures using Reinforcement Learning. In *Proceedings of International Conference on Reinforcement Learning 2017*.
- [5] Krizhevsky, A & Hinton, G. 2009. *Learning multiple layers of features from tiny images - Tech. Rep. 1*. Computer Science Department. University of Toronto, Toronto Canada.
- [6] Yuval Netzer, Tao Wang, Adam Coates, Alessandro Bissacco, Bo Wu, Andrew Y. Ng. 2011. Reading Digits in Natural Images with Unsupervised Feature Learning. In *NIPS Workshop on Deep Learning and Unsupervised Feature Learning 2011*.
- [7] Esteban Real, Sherry Moore, Andrew Selle, Saurabh Saxena, Yutaka Leon Suematsu, Jie Tan, Quoc Le, Alex Kurakin. Large-Scale Evolution of Image Classifiers. In *Proceedings of 34th International Conference on Machine Learning 2017*.
- [8] Sagar Sharma. Sept. 6, 2017. *Activation Functions in Neural Networks Towards Data Science*. Medium.



Questions?